

Diffusing ideas

Adventures with noise and building mathematical toys

Robert J. Hardwick

January 5, 2023

Introduction

Diffusing ideas records a journey of research exploration and software development over several years. It's primarily an expression of interrelated ideas around the statistical inference and automated control of stochastic phenomena in general. However, in order to manifest these ideas into reality, this project has also involved designing and building a lot of new open-source scientific software written in the [Python](#) and [Go](#) programming languages as well. The main motivation behind developing these computational elements is to form a nice foundation from which anyone can build loads of new applications, and the hope is that an additional benefit should be that they make scientific research and discovery of new phenomena a more efficient experience for the scientist.

Having just stated all this 'altruistic' stuff, I would be remiss if I left the motivations for writing this book at that. The necessity for testing the software has provided a great excuse to study, play with and invent new derivatives of an extensive range of toy mathematical models and imagining some of these situations where the code might be useful has been a lot of fun. I'm also not going to write an 'about the author' section because I think they are usually quite boring. Suffice it to say that I have a background in various fields of science which I hope will make the mathematical model digressions quite varied and interesting to the reader!

A quick note on the software; any software that I describe in this book (including the [software which compiles the book itself](#)) will always be shared under a [MIT License](#) in a public Git repository.¹ Forking these repositories and submitting pull requests for new features or applications is strongly encouraged too, though I apologise in advance if I don't follow these up very quickly as all of this work has to be conducted independently in my free time outside of work hours.

No quest would be complete without a guide, so in this introductory section, I think it makes the most sense to now outline some of the key milestones of the book and their motivations within the context of the overall research project. My core aims can be divided into proposing answers to a set of interdependent research questions, in the following order:

1. How do we simulate a general set of stochastic phenomena?
2. How do we then learn/identify 1. from real-world data?
3. How do we simulate a general set of automated control policies to interact with 1.?
4. How do we then optimise 3. to achieve a specified control objective?

¹The repositories will always be somewhere on this list: <https://github.com/umbralcalc?tab=repositories>.

Table of contents

1	Building a generalised simulator	1
1.1	Core mathematical formalism	1
1.2	Flavours of noise with continuous sample paths	2
1.3	Flavours of noise with discontinuous sample paths	2
1.4	Summary of desirable features	3
1.5	Software design choices	4
2	Simulating a financial market	5
2.1	Introducing Q-Hawkes processes	5
3	Quantum jumps on generic networks	7
3.1	The Lindblad equation	7
4	Inferring dynamical 2D maps	9
4.1	Adapting the stochadex formalism	9
5	Learning from ants on curved surfaces	11
5.1	Diffusive limits for ant interactions	11
6	Hydrodynamic ensembles from input data	13
6.1	The Boltzmann/Navier-Stokes equations	13
7	Generalised statistical inference tools	15
7.1	Likelihood-free methods	15
8	Interacting with systems in general	17
8.1	Parameterising general interactions	17
9	Managing a Rugby match	19
9.1	Introduction	19
9.2	Designing the event simulation engine	19
9.3	Linking to player attributes	20
9.4	Deciding on gameplay actions	20
9.5	Writing the game itself	20

10 Influencing house prices	23
11 Optimising actions for control objectives	25
12 Resource allocation for epidemics	27
13 Quantum system control	29
14 Abstract models	31

Building a generalised simulator

Concept. To design and build a generalised simulation engine that is able to generate samples from a ‘Pokédex’ of possible stochastic processes that a researcher might encounter. A ‘Pokédex’ here is just my fanciful description for a very general class of multidimensional stochastic processes that pop up everywhere in taming the mathematical wilds of real-world phenomena, and which also leads to a name for the software itself: the ‘stochadex’. With such a thing pre-built and self-contained, it can become the basis upon which to build generalised software solutions for a lot of different problems. For the mathematically-inclined, this chapter will require the introduction of a new formalism which we shall refer back to throughout the book. For the programmers, the public Git repository for the code that is described in this chapter can be found here: <https://github.com/umbralcalc/stochadex>.

1.1 Core mathematical formalism

Before we dive into the software, we need to define the mathematical approach that we’re going to take in order to be able to describe a really general set of stochastic phenomena. From experience, it seems reasonable to start by writing down the following formula which describes iterating some arbitrary process forward in time (by one finite step) and adding a new row each to some matrices $V' \rightarrow V$ and $X' \rightarrow X$

$$X_{t+1}^i = F_{t+1}^i(X', V', t) \quad (1.1)$$

$$V_{t+1}^i = X_{t+1}^i - X_t^i, \quad (1.2)$$

where: i is an index for the dimensions of the ‘state’ space; t is the current time index for either a discrete-time process or some discrete approximation to a continuous-time process; X is the next version of X' after one timestep (and hence one new row has been added); V as the next version of V' (similarly to X and X'); and $F_{t+1}^i(X', V', t)$ as the latest element of an arbitrary matrix-valued function.

So the basic computational idea here is to iterate the matrices X and V forward in time by a row, and use their previous versions (X' and V') as entire matrix inputs into functions which populate

the elements of their latest rows. But why go to all this trouble of storing matrix inputs for previous values of the same process? For a large class of stochastic processes this memory of past values is essential to consistently construct the sample paths moving forward. This is true in particular for *non-Markovian* phenomena, where the latest values don't just depend on the immediately previous ones but can depend on values which occurred much earlier in the process.

Note that, based on the definition in Eq. (1.2) above, the following relation is also valid

$$X_t^i = X_s^i + \sum_{s'=s+1}^t V_{s'}^i, \quad (1.3)$$

where $s < t$.

1.2 Flavours of noise with continuous sample paths

For *Wiener process noise*, adopting the [Itô interpretation](#) in this section, we can define W_t^i is a sample from a Wiener process for each of the state dimensions indexed by i and our formalism becomes

$$X_{t+\delta t}^i = F_{t+\delta t}^i(X', t) + \textcolor{red}{W}_{t+\delta t}^i - W_t^i. \quad (1.4)$$

Other interpretations of the noise are less immediately compatible with our formalism as it is currently written, e.g., [Stratonovich](#) or others within the α -family, but it seems less necessary to complicate the details of this section further, so we'll just cover these extensions at the software implementation level. Note also that we may also allow for correlations between the noises in different dimensions.

For *Geometric Brownian motion noise*, we simply have

$$X_{t+\delta t}^i = F_{t+\delta t}^i(X', t) + \textcolor{red}{X}_t^i(\textcolor{red}{W}_{t+\delta t}^i - W_t^i). \quad (1.5)$$

And say, e.g., *fractional Brownian motion noise*, where $B_t^i(H_i)$ is a sample from a fractional Brownian motion process with Hurst exponent H_i for each of the state dimensions indexed by i , we simply substitute again

$$X_{t+\delta t}^i = F_{t+\delta t}^i(X', t) + \textcolor{red}{B}_{t+\delta t}^i(H_i) - B_t^i(H_i). \quad (1.6)$$

Generalised continuous noises would take the form

$$X_{t+\delta t}^i = F_{t+\delta t}^i(X', t) + \textcolor{red}{g}_{t+\delta t}^i(X', \textcolor{red}{W}_{t+\delta t}^i - W_t^i, \dots), \quad (1.7)$$

where $g_{t+\delta t}^i(X', \textcolor{red}{W}_{t+\delta t}^i - W_t^i, \dots)$ is some continuous function of its arguments which can be expanded out with [Itô's Lemma](#).

1.3 Flavours of noise with discontinuous sample paths

Jump process noises generally could take the form

$$X_{t+\delta t}^i = F_{t+\delta t}^i(X', t) + \textcolor{red}{J}_{t+\delta t}^i(X', \dots), \quad (1.8)$$

where $J_{t+\delta t}^i(X', \dots)$ are samples from some arbitrary jump process (e.g., compound Poisson) which could generally depend on a variety of inputs, including X' .

Poisson process noises would generally take the form

$$X_{t+\delta t}^i = F_{t+\delta t}^i(X', t) + N_{t+\delta t}^i(\lambda_i) - N_t^i(\lambda_i), \quad (1.9)$$

where $N_t^i(\lambda_i)$ is a sample from a Poisson process with rate λ_i for each of the state dimensions indexed by i . Note that we may also allow for correlations between the noises in different dimensions.

Time-inhomogeneous Poisson process noises would generally take the form

$$X_{t+\delta t}^i = F_{t+\delta t}^i(X', t) + N_{t+\delta t}^i(\lambda_{t+\delta t}^i) - N_t^i(\lambda_t^i), \quad (1.10)$$

where λ_t^i is a deterministically-varying rate for each of the state dimensions indexed by i .

Cox (doubly-stochastic) process noises would generally take the form

$$X_{t+\delta t}^i = F_{t+\delta t}^i(X', t) + N_{t+\delta t}^i(\Lambda_{t+\delta t}^i) - N_t^i(\Lambda_t^i), \quad (1.11)$$

where the rate Λ_t^i is now a sample from some continuous-time stochastic process (in the positive-only domain) for each of the state dimensions indexed by i .

Self-exciting process noises would generally take the form

$$X_{t+\delta t}^i = F_{t+\delta t}^i(X', t) + N_{t+\delta t}^i[\mathcal{I}_{t+\delta t}^i(N', \dots)] - N_t^i[\mathcal{I}_t^i(N'', \dots)], \quad (1.12)$$

where the stochastic rate $\mathcal{I}_t^i(N', \dots)$ now depends on the history of N' explicitly (amongst other potential inputs - see, e.g., [Hawkes processes](#)) for each of the state dimensions indexed by i .

Generalised probabilistic discrete state transitions would take the form

$$X_{t+\delta t}^i = F_{t+\delta t}^i(X', t) + T_{t+\delta t}^i(X'), \quad (1.13)$$

where $T_{t+\delta t}^i(X')$ is a generator of the next state to occupy. This generator uses the current state transition probabilities (which are generally conditional on X') at each new step.

1.4 Summary of desirable features

- using the learnings from the previous sections looking at specific example processes
- above formalism is so general that it can do anything - so while it shall serve as a useful guide and reference point, it would be good here to go through more of the specific desirable components we want to have access to in the software itself
- it might not always be convenient to have the windowed histories stored as S but some other varying quantity which can be used to construct S? take fractional brownian motion as an example of this! hence, need to provide more possible input histories into S
- want the timestep to have either exponentially-sampled lengths or fixed lengths in time
- formalism already isn't explicit about the choice of deterministic integrator in time
- but also want to be able to choose the stochastic integrator in continuous processes (Itô or Stratonovich?)

- enable correlated noise terms at the sample generator level
- configurable setup of simulations with just yamls + a single .go file defining the terms

Test cite [\[1\]](#)

1.5 Software design choices

Ideally, the stochadex sampler should be designed to try and maintain a balance between performance and flexibility of utilisation.

Simulating a financial market

Concept. The idea here is to use the Q-Hawkes processes and the Bouchaud work to come up with some interesting simulations of financial markets.

2.1 Introducing Q-Hawkes processes

Quantum jumps on generic networks

Concept. The idea is to follow this sort of thing [here](#) to simulate the Lindblad equation over an arbitrary network of entangled states.

3.1 The Lindblad equation

Inferring dynamical 2D maps

Concept. The idea here is

4.1 Adapting the stochadex formalism

Learning from ants on curved surfaces

Concept. The idea here is

5.1 Diffusive limits for ant interactions

Hydrodynamic ensembles from input data

Concept. The idea here is

6.1 The Boltzmann/Navier-Stokes equations

Generalised statistical inference tools

Concept. The idea here is to extend the stochadex with tools for very generalised statistical inference (ABC algorithms and the like) that will work in nearly every situation. Probably need to exploit the phase space analogy of the formalism.

7.1 Likelihood-free methods

Interacting with systems in general

Concept. The idea here is

8.1 Parameterising general interactions

Managing a Rugby match

Concept. The idea here is

9.1 Introduction

Since the basic game engine will run using the [stochadex](#) sampler, the novelties in this project are all in the design of the rugby match model itself. And, in this instance, I'm not especially keen on spending a lot of time doing detailed data analysis to come up with the most realistic values for the parameters that are dreamed up here. Even though this would also be interesting.

One could do this data analysis, for instance, by scraping player-level performance data from one of the excellent websites that collect live commentary data such as [rugbypass.com](#) or [espn.co.uk/rugby](#).

This game is primarily a way of testing out the interface of the stochadex for other users to build projects with. This should help to both iron out some of the kinks in the design, as well as prioritise adding some more convenience methods for event-based modelling into its code base.

9.2 Designing the event simulation engine

We need to begin by specifying an appropriate event space to live in when simulating a rugby match. It is important at this level that events are defined in quite broadly applicable terms, as it will define the state space available to our stochastic sampler and hence the simulated game will never be allowed to exist outside of it. So, in order to capture the fully detailed range of events that are possible in a real-world match, we will need to be a little imaginative in how we define certain gameplay elements when we move through the space.

The diagrams below sum up what should hopefully work as a decent initial approximation while providing a little context with specific examples of play action.




Figure 9.1: Simplified event graph of a rugby union match.

9.3 Linking to player attributes

9.4 Deciding on gameplay actions

9.5 Writing the game itself

	Run Phase	Kick Phase	Ruck	Maul	Lineout	Scrum	Penalty	Free Kick	Distribution Model
Ball Win/Lose Binomial	✓	✓	✓	✓	✓	✓			Binomial Probability
Concede Penalty / Free Kick Binomial	✓	✓	✓	✓	✓	✓			Binomial Probability
Distance Gained / Lost Gamma	✓	✓	✓	✓		✓			Gamma Distribution
Knock-on Binomial	✓	✓							Binomial Probability
Try Binomial	✓	✓					✓		Binomial Probability
Kicking Normal		✓					✓	✓	Normal Distribution

Figure 9.2: Optional model ideas.

10

Influencing house prices

Concept. The idea here is

Optimising actions for control objectives

Concept. The idea controlhere is

Resource allocation for epidemics

Concept. The idea here is to limit the spread of some abstract epidemic through the correct time-dependent resource allocation.

Quantum system control

Concept. The idea here is to follow stuff along these lines [here](#).

Abstract models

Concept. The idea here is

Bibliography

Bibliography

- [1] E. Dimastrogiovanni, M. Fasiello, R. J. Hardwick, H. Assadullahi, K. Koyama and D. Wands, *Non-Gaussianity from Axion-Gauge Fields Interactions during Inflation*, [1806.05474](#).